

CSC 3150 Project 3 Report

Zhu Chuyan 119010486 at *The Chinese University of Hong Kong, Shenzhen*

Table of contents

- [CSC 3150 Project 3 Report](#)
 - [Table of contents](#)
 - [1. How to run the program](#)
 - [1.1 Running environments](#)
 - [1.2 Executions](#)
 - [Main task: Single Process VM Simulation](#)
 - [Bonus: Multiple Process VM Simulation](#)
 - [2. Implementations](#)
 - [2.1 Single Process VM Simulation](#)
 - [Inverted Tables](#)
 - [Table Searching](#)
 - [LRU Policy Implementations](#)
 - [Memory Layouts](#)
 - [2.2 Multiple Process VM Simulation](#)
 - [3. The results](#)
 - [Single Process VM Simulation](#)
 - [My User Program](#)
 - [Multiple Processes VM Simulation \(bonus\)](#)
 - [4. What I learn from the project](#)

1. How to run the program

1.1 Running environments

- Operating system: `Linux`
- Linux distribution: `CentOS Linux release 7.6.1810 (Core)`
- Linux kernel version: `Linux version 3.10.0-957.21.3.el7.x86_64`
- Compiler/CUDA version: `Cuda compilation tools, release 10.1, V10.1.105`

1.2 Executions

Main task: Single Process VM Simulation

A Makefile has been prepared, so you can compile with the following command:

```
cd path_to_project/source/  
make
```

You can then test the program by executing

```
./vm
```

Bonus: Multiple Process VM Simulation

```
cd path_to_project/bonus/  
make  
# if you wish to clean the object files:  
make clean
```

Then you can simply run by:

```
./vm_bonus
```

2. Implementations

2.1 Single Process VM Simulation

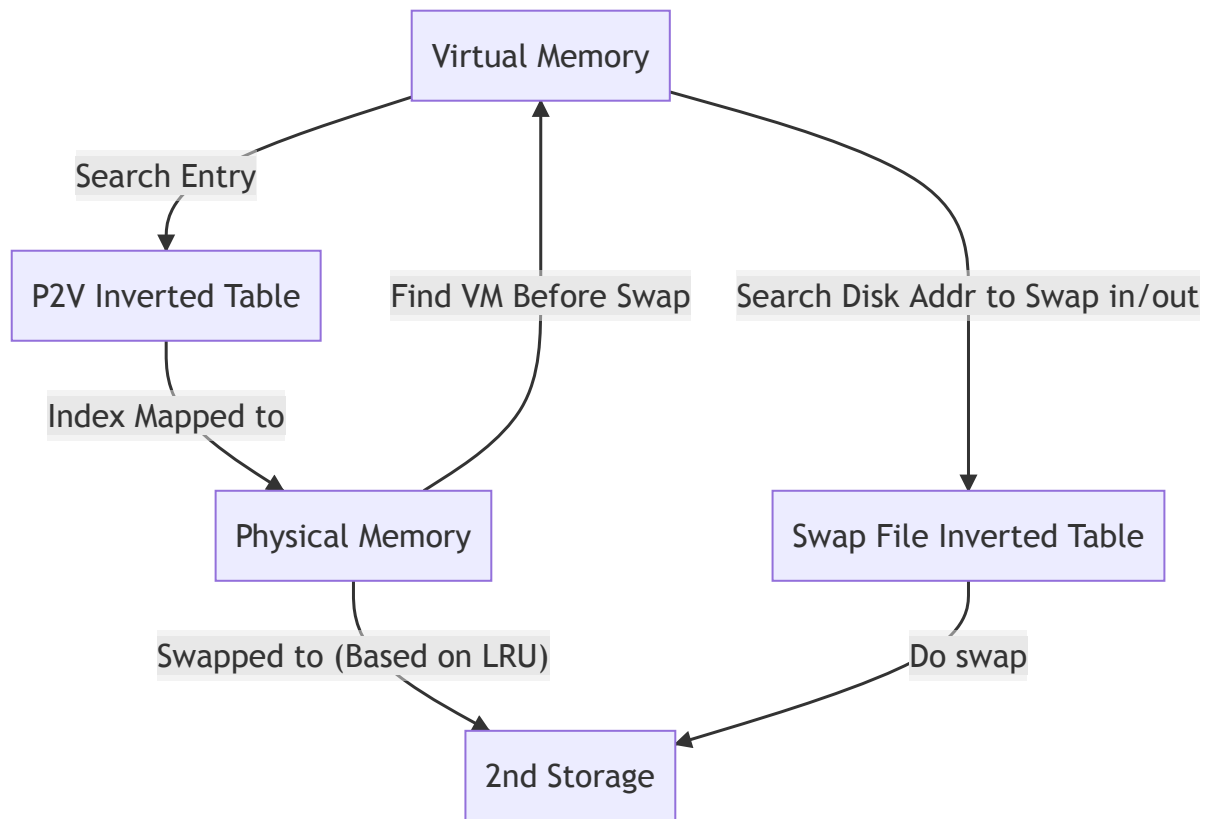


Figure 1: Main Program Implementation Flow Chart

Inverted Tables

First of all, we need a page table for physical memory-virtual memory mapping. Due to the limitation of memory space, an inverted table is needed to implement.

The entry of the table is the virtual memory address (page address, page offset excluded), and the index of the table is the physical memory address, therefore, there will be in total 1024 entries as there are 1024 pages in the physical memory according to the given assumption.

As the upper limit of the VM address space is 256K, larger than the secondary storage size, therefore, I did not adopt the one-to-one mapping between the swap file location and the VM address, instead, another inverted table is introduced to record the relationship.

Table Searching

As yet another time-space tradeoff, the inverted table requires us to search for the specific VM address, and it may takes time complexity of $O(n)$, which is quite inefficient. Therefore, a hashing method may help to improve the performance.

However, the hashing result may generate outputs confliting with the LRU policy when the page table is full, therefore, a linear search will be adopted in my program **after the table is full**. Before that, a open-addressing hashing method is used. All entries are firstly marked as invalid, and the linear probing process will terminate whenever it encounters a invalid entry or finds a desired entry.

LRU Policy Implementations

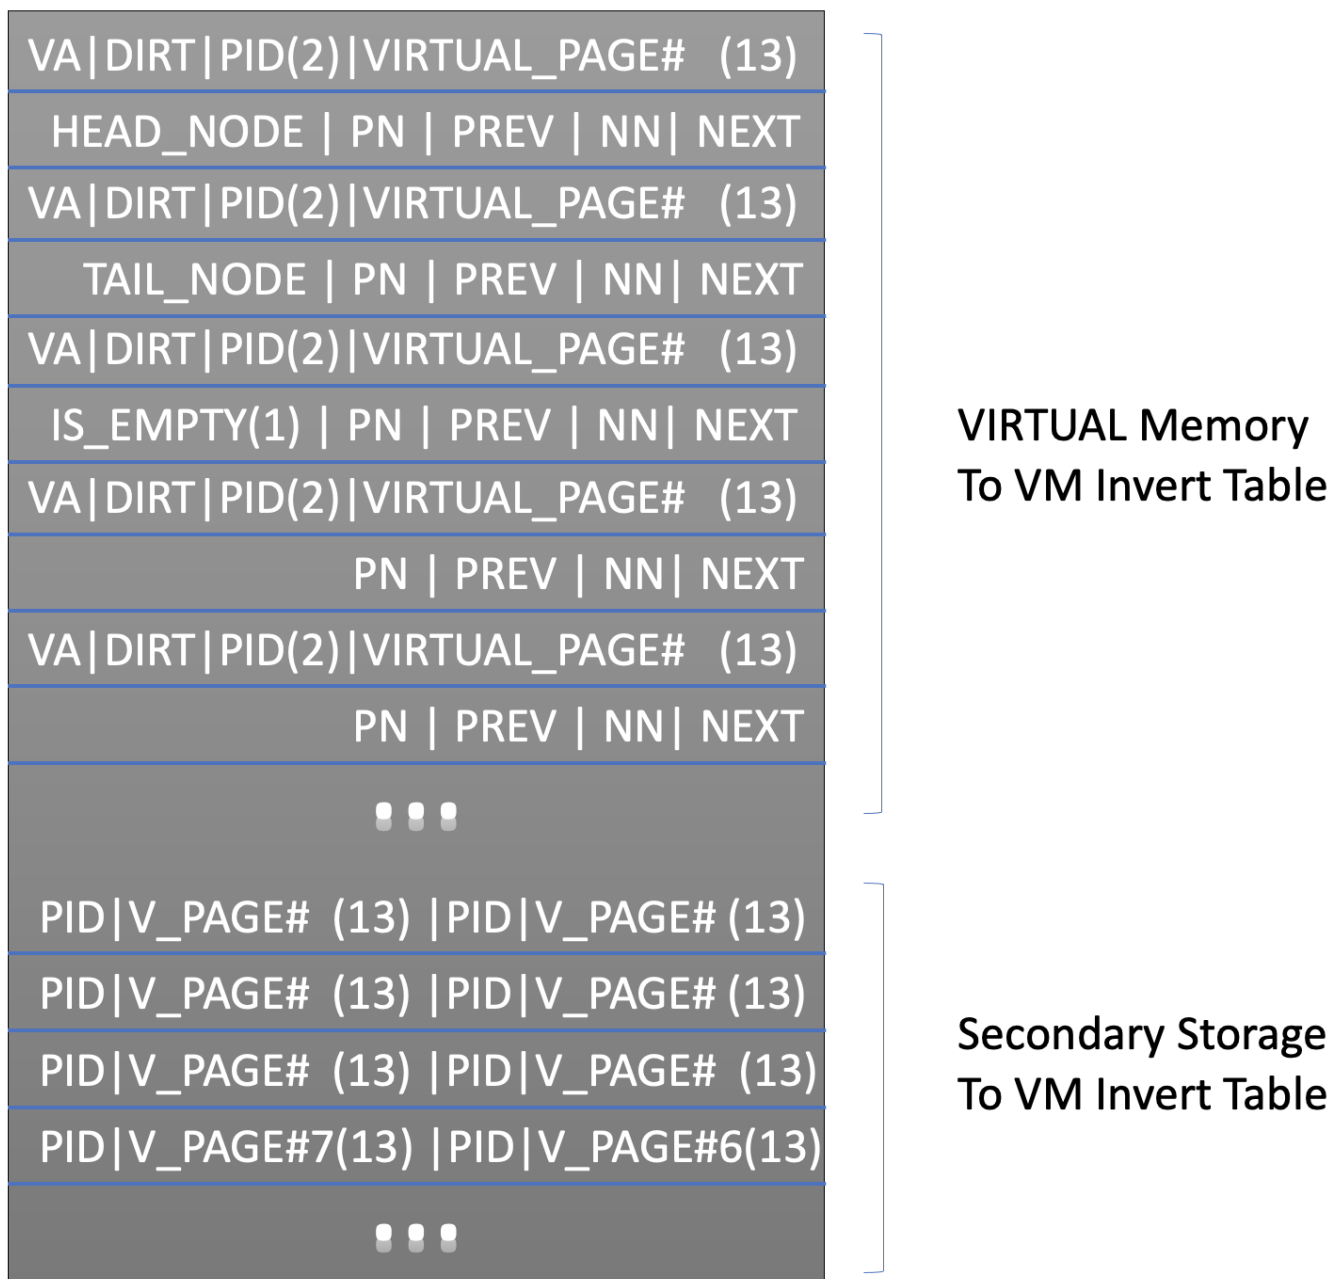
I implemented LRU with a doubly linked list to reduce the searching time, for each entry in the inverted page table, there is a pointer pointing to the previous and next entries in term of latest visiting time (like a priority queue).

Each time a memory address (physical) is accessed, the corresponding node will be moved to the head of the linked list, a each time a victim is to be selected to evict, it will be the tail of the node.

Memory Layouts

As only 16KB memory is available for all tables and related records, the memory is quite compacted to fulfills the designing ideas mentioned above, therefore, the memory layout is shown as below and dozens of macros are written to facilitate the use.

Memory Layouts



The components are elaborated as follows:

- VA: Valid bit
- DIRT: Dirty bit
- PID: the pid of the program (applicable in bonus)
- Virtual Page: the corresponding virtual page number
- Head_node: the head node of the linked list
- Tail_node: the tail node of the linked list
- PREV: Previous node in the linked list
- NEXT: Next node in the linked list
- PN: The previous node is null
- NN: The next node is null

Note that in the Physical-memory-to-VM inverted table, each entry occupies two 32-bit block and in the Secondary-storage-to-VM inverted table, each entry occupies half 32-bit block.

2.2 Multiple Process VM Simulation

The multiple process simulation is done with multiple CUDA threads within a same block. It is hard to achieve the non-preemptive priority scheduling as it is almost impossible to implement the lock for the CUDA threads within **the same block** due to its wrap architecture. Therefore, the multiple processes rather run sequentially to simulate the non-preemptive priority scheduling to avoid the data complict in accessing the same inverted page table.

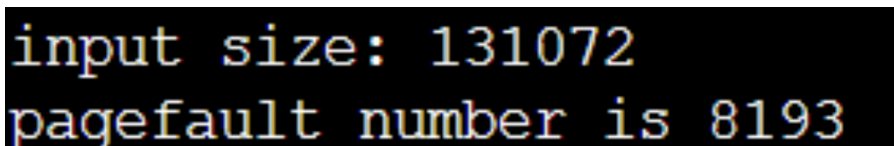
The sequantial execution is achived by leveraging the CUDA `__syncthreads()` and a if condition.

What's more, the pid is added to the head of the VM address to distinguish the addresses of different processes, and other implementation details stay consistent with the main part mentioned above.

3. The results

Single Process VM Simulation

- Page fault number: 8093
- Screenshots of result:



```
input size: 131072
pagefault number is 8193
```

Figure 1: Single Process Test Results

- Analysis of the results:
 - First of all, we invoke 4k page write to the file, causing 4k page faults
 - Then the following read function calls from 4k down will causes exactly 1 pagefaults.
 - Finally the `vm_read()` calls in the snapshot will further generates 4k pagefaults.

My User Program

For required work:

1

```
for (int i = 0; i < input_size; i++)
    vm_write(vm, i, input[i]);

for (int i = input_size - 1; i >= input_size - 32769; i--)
    int value = vm_read(vm, i);

vm_snapshot(vm, results, 0, input_size);
```

PF: 8193

2

```
for (int i = 0; i < input_size; i++)
    if(i%2) vm_write(vm, i, input[i]);

for (int i = input_size - 1; i >= input_size - 32769; i--)
    if(i%2) int value = vm_read(vm, i);

for (int i = 0; i < input_size; i++)
    if(i%2==0) vm_write(vm, i, input[i]);

vm_snapshot(vm, results, 0, input_size);
```

PF: 12289

Multiple Processes VM Simulation (bonus)

- Using the user program I provides, 32772 page faults are generated.

```

for (int i = 0; i < input_size; i++){
    if(i%4==pid){
        vm_write(vm, i, input[i],pid);
    }
}

for (int i = input_size - 1; i >= input_size - 32769; i--)
    if(i%4==pid) int value = vm_read(vm, i,pid);

for(u32 i = 0; i < input_size; i++) {
    if(i%4==pid) results[i] = vm_read(vm,i,pid);
}
//No snapshot is further needed

```

4. What I learn from the project

I learn the principles of paging memory management system and how to implement them. I also learn that the replacement algorithm (for example, LRU we use in the project) is of great importance in a paging system, and the utilize the memory as much as possible.